

Sparse Matrices Correctness 2019 Paper

ACM Reference Format:

. 2019. Sparse Matrices Correctness 2019 Paper. In *Proceedings of ACM Conference (Conference'17)*. ACM, New York, NY, USA, ?? pages. <https://doi.org/10.1145/nnnnnnn.nnnnnnn>

1 INTRODUCTION

Sparse matrix data formats can be used to compress large matrices with a small number of non-zero elements into a more efficient representation. Scientific and engineering software that make use of sparse matrix formats are often implemented in low-level imperative languages such as C++ and Fortran. The optimized nature of these software often means that the structural organization of sparse matrix formats and mathematical computations are heavily intertwined. Additionally, the myriad data formats and complexities involved in tuning the operations on these formats to achieve efficient implementations means the development of software that utilizes sparse matrices can be a tedious and error-prone task.

Often, the solvers provided by popular linear algebra libraries used in scientific computing require assembled sparse matrices as input, leaving the entire assembly process to the developer. Considering the problem domains in which they are used, this can quickly become a complex task. For example, sparse matrices are often used to represent the meshes used in the finite element and finite difference methods. In order to represent physical properties, these meshes can contain rich datasets, the values of which may depend on the complex spatial relationships of the mesh components. As a result, the matrix assembly process can quickly become intertwined with the mathematics required to determine the precise values that will be placed in the matrix. Furthermore, as performance improvements are made by, for example, exploiting parallelism in modern hardware, the assembly process and embedded calculations can become even more complex.

To address the difficulties in assembling a sparse matrix, object-oriented libraries such as PETSc [?] and Eigen [?] provide data abstractions targeted towards specific classes of solvers. These libraries allow developers to assemble sparse matrices without having to address the structural complexities that a sparse matrix format may present. However, the utility of these libraries may be limited due to solver limitations or other performance restrictions.

Alternatively, compiler support may help with the development and performance tuning of software that makes use of sparse matrices. Some compilers [?] allow developers to work with dense matrices in code, generating a sparse program at compile time. Others [?] allow the user to provide the compiler with an abstract description of a sparse format, from which the compiler can make automatic optimizations in code.

To mitigate the introduction of errors into code during development, techniques such as unit testing and test-driven development are often employed. These methods aim to reveal bugs by testing the individual components of a piece of software during development. While useful, these methods are both tedious, requiring test cases to be written manually, and incomplete, meaning they are incapable of checking every scenario or combination of scenarios.

We describe an approach that allows developers to reason about the inherent complexities of sparse matrix formats and operations without limiting implementation choices. Elements of this approach include declarative modeling and automatic, push-button analysis using the Alloy Analyzer [?], a lightweight bounded model checking tool. Characteristics of sparse matrix formats and operations are modeled using abstraction based methods [?], including data [?] and predicate [?] abstraction, data and functional refinement [?], and other techniques, manually, as part of a modeling process.

The benefits of this approach lie in its generality. In contrast to code, Alloy allows for partial descriptions in which details are hidden behind abstractions. This allows us to, for example, reason about the complexities of building a sparse matrix without the burden of calculating the specific values that matrix will hold. Furthermore, in being able to freely choose the level of abstraction, we are able to model a wide range of —, from specific algorithms written in low-level imperative languages to higher level, abstract concepts.

- partial descriptions - level of abstraction - structure and behavior

The remainder of this paper is organized as follows. In Section ?? we discuss related work. In Section ?? we describe certain elements of lightweight formal methods and how they will be applied to our problem domain using Alloy. In Section ?? we model three matrix formats, including one abstract matrix and two sparse matrix formats. In Section ?? we model the translation of a sparse matrix from the ELL format to the CSR format, and in Section ?? we model a sparse matrix-vector multiplication algorithm for the CSR format. Finally, in Section ?? we conclude with future directions and closing remarks.

2 RELATED WORK

Arnold et al. [?] design a functional language and proof method for the automatic verification of sparse matrix codes.

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than ACM must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

Conference'17, July 2017, Washington, DC, USA

© 2019 Association for Computing Machinery.

ACM ISBN 978-x-xxxx-xxxx-x/YY/MM. . . \$15.00

<https://doi.org/10.1145/nnnnnnn.nnnnnnn>

Their “little language” (LL) can be used to specify sparse codes as functional programs in which computations are sequences of high-level transformations on lists. These models are then automatically translated into Isabelle/HOL where they can be verified for full functional correctness. The authors use this automated proof method to verify a number of sparse matrix formats and their sparse matrix-vector multiplication operations. Their approach is purely functional, relying on typed λ -calculus and Isabelle libraries to perform proofs using Isabelle/HOL. LL, a shallow embedding in Isabelle/HOL, provides simple, composable rules that can be used to fully describe a sparse matrix format, removing the burden of directly writing proofs. In quantifying rule reuse, the authors find that “on average, fewer than 19% of rules used by a particular format are specific to this format, while over 66% of these rules are used by at least three additional formats, a significant level of reuse. Even of the rules needed for more complex formats (CSC and JAD), only up to a third are format-specific.” [?] The method enables a significant amount of rule reuse, but does not entirely prevent one from having to write some amount of λ -calculus. Filling in these gaps may prove difficult for those without a strong background in functional programming and theorem proving. Our approach, on the other hand, may appeal to an audience of scientists and engineers, who are accustomed to working with models anyway, and with the kind of automatic, push-button analysis supported by Alloy, those who develop software can focus on modeling and design instead of theorem proving.

Arnold approach does not support modeling of assembly of sparse matrices. Relational approach in combination with state change allows us to model the assembly process as well as updates to matrix values. Relational approach also allows for modeling of format conversions.

REFERENCES

- [1] S. Balay, S. Abhyankar, M. F. Adams, J. Brown, P. Brune, K. Buschelman, L. Dalcin, A. Dener, VictorEijkhout, W. D. Gropp, D. Karpeyev, D. Kaushik, M. G. Knepley, D. A. May, L. C. McInnes, R. T. Mills, T. Munson, K. Rupp, P. Sanan, B. F. Smith, S. Zampini, H. Zhang, and H. Zhang, “PETSc Web page,” <https://www.mcs.anl.gov/petsc>, 2019. [Online]. Available: <https://www.mcs.anl.gov/petsc>
- [2] G. Guennebaud, B. Jacob *et al.*, “Eigen v3,” <http://eigen.tuxfamily.org>, 2010.
- [3] A. Bik and H. Wijshoff, “Advanced compiler optimizations for sparse computations,” *Journal of Parallel and Distributed Computing*, vol. 31, no. 1, pp. 14 – 24, 1995. [Online]. Available: <http://www.sciencedirect.com/science/article/pii/S0743731585711410>
- [4] A. J. C. Bik and H. A. G. Wijshoff, “Automatic data structure selection and transformation for sparse matrix computations,” *IEEE Transactions on Parallel and Distributed Systems*, vol. 7, no. 2, pp. 109–126, Feb 1996.
- [5] V. Kotlyar, K. Pingali, and P. Stodghill, “A relational approach to the compilation of sparse matrix programs,” in *Euro-Par’97 Parallel Processing*, C. Lengauer, M. Griebl, and S. Gortlatch, Eds. Berlin, Heidelberg: Springer Berlin Heidelberg, 1997, pp. 318–327.
- [6] D. Jackson, *Software Abstractions: Logic, Language, and Analysis*. The MIT Press, 2012.
- [7] E. M. Clarke, O. Grumberg, and D. E. Long, “Model checking and abstraction,” *ACM Transactions on Programming Languages and Systems (TOPLAS)*, vol. 16, no. 5, pp. 1512–1542, 1994.
- [8] J. Dingel and T. Filkorn, “Model checking for infinite state systems using data abstraction, assumption-commitment style reasoning and theorem proving,” in *Computer Aided Verification: 7th International Conference, CAV ’95 Liège, Belgium, July 3–5, 1995 Proceedings*. Berlin, Heidelberg: Springer Berlin Heidelberg, 1995, pp. 54–69. [Online]. Available: https://doi.org/10.1007/3-540-60045-0_40
- [9] S. Graf and H. Saïdi, “Construction of abstract state graphs with PVS,” in *International Conference on Computer Aided Verification*. Springer, 1997, pp. 72–83.
- [10] J. Woodcock, *Using Z : specification, refinement, and proof*. London New York: Prentice Hall, 1996.
- [11] G. Arnold, J. Hlzl, A. Sinan Kksal, R. Bodik, and M. Sagiv, “Specifying and verifying sparse matrix codes,” *ACM SIGPLAN Notices*, vol. 45, pp. 249–260, 09 2010.
- [12] E. Jones, T. Oliphant, P. Peterson *et al.*, “SciPy: Open source scientific tools for Python,” 2001–. [Online]. Available: <http://www.scipy.org/>
- [13] J. Demmel, J. Dongarra, A. Ruhe, and H. van der Vorst, *Templates for the Solution of Algebraic Eigenvalue Problems: A Practical Guide*, Z. Bai, Ed. Philadelphia, PA, USA: Society for Industrial and Applied Mathematics, 2000.
- [14] Y. Saad, “Sparskit: a basic tool kit for sparse matrix computations - version 2,” 1994.
- [15] A. Milicevic, J. P. Near, E. Kang, and D. Jackson, “Alloy*: A general-purpose higher-order relational constraint solver,” in *2015 IEEE/ACM 37th IEEE International Conference on Software Engineering*, vol. 1, May 2015, pp. 609–619.
- [16] T. Dyer, “Article github repository,” 2019. [Online]. Available: <https://github.com/atdyer/alloy-lib>